

Lattice QCD GPU programming with Quda and Grid

November 7, 2025

Outline

- ▶ slaph interface in QUDA
 - ▶ meson functions
 - ▶ Baryon kernel

Recap on the last lecture

We learned that there is an interface for *GEMM routines in QUDA and it will translate automatically between row-major and column-major formats. The interface is shallow and it will be needed for the rest of this lecture.

There is no standard batched BLAS function call - MKL differs from cuBLAS and HipBLAS although the latter agree.

We introduced compilation and running with GRID, we talked a bit about the structure of GRID and how that translates to effective parallelism.

We saw that GRID hides the GPU interface from us but we can (and often should) drop down and use these "accelerator_for" loops, as substantial speedup is possible.

Meson currents

I showed the inner loop of the meson contractions in the last lecture, and teased it in the second. To recap we want to implement N DFTs as a small matrix batched with several right-hand-side larger matrices.

$$N_P \left\{ \begin{bmatrix} e^{ip_0 \cdot x_0} & \dots & e^{ip_0 \cdot x_{L^3}} \\ \vdots & \vdots & \vdots \\ e^{ip_{N_P-1} \cdot x_0} & \dots & e^{ip_{N_P-1} \cdot x_{L^3}} \end{bmatrix} \right. \overbrace{\quad}^{L^3} \times L^3 \left\{ \begin{bmatrix} Tr(x_0, t_0) & \dots & Tr(x_0, t_{L_t-1}) \\ \vdots & \vdots & \vdots \\ Tr(x_{L^3}, t_0) & \dots & Tr(x_{L^3}, t_{L_t-1}) \end{bmatrix} \right. \overbrace{\quad}^T$$

Where the "N" is our "blockSizeMomProj" - a division of the $n1 \times n2$ noise possibilities we will run through.

What is the code doing? Well it is performing a DFT over traces of meson noise combinations $n1$ and $n2$

$$C(n1, n2, p, t) = \sum_x e^{ip \cdot x} \text{Tr}_C [\psi_{n2}(x, t) \bar{\psi}_{n1}(0)]$$

Where Tr_C is a trace over color indices.

Note that $n1$ moves slowest in this construction and t fastest. The trace is just the kernel we created in lecture 2.

As this is an interface we expect the "latticeColorVectors" to be on the host. The function we will call is:

```
void laphCurrentKernel( const int n1, const int n2,
                        const int nMom,
                        const int blockSizeMomProj,
                        void **host_quark,
                        void **host_quark_bar,
                        const double _Complex *host_mom,
                        QudaInvertParam inv_param,
                        void *return_array,
                        const int X[4])
```

All arguments here exist on the host. Here "host_mom" is a flattened array of $Np \times L^3$ Fourier twiddle factors. $X[4]$ carries the local dimensions.

Allocations

Here we allocate memory on the device, skipping creating the CPU colorSpinorField creation

```
const size_t data_ret_bytes = nMom*X[3]*n1*n2*2*precision;
const size_t data_tmp_bytes = n_sites*2*precision*blockSizeMom;
const size_t data_mom_bytes = nMom*n_spatial_sites*2*precision;
void *d_ret = pool_device_malloc(data_ret_bytes);
void *d_tmp = pool_device_malloc(data_tmp_bytes);
void *d_mom = pool_device_malloc(data_mom_bytes);
```

Here we need to specify the return array size for the device "d_ret" a temporary that holds the lattice-wide traces "d_tmp" and some space for the copied momentum twiddles "d_mom"

batched-strided ZGEMM parameters

We call the constructor

```
QudaBLASParam cublas_param_mom_sum = newQudaBLASParam();
```

And specify that the matrix to the right is transposed as our fields are in c-order

```
cublas_param_mom_sum.trans_a = QUDA_BLAS_OP_N;
cublas_param_mom_sum.trans_b = QUDA_BLAS_OP_T;
cublas_param_mom_sum.m = nMom ;
cublas_param_mom_sum.n = X[3] ;
cublas_param_mom_sum.k    = n_spatial_sites ;
cublas_param_mom_sum.lda = n_spatial_sites ;
cublas_param_mom_sum.ldb = n_spatial_sites ;
cublas_param_mom_sum.ldc = X[3] ;
```

"k" is our inner index for the matrix multiply and we will be creating an $nMom \times X[3]$ output matrix.

BLAS params continued

Here we are saying that there will be no striding (i.e. we reuse the same matrix) in the left matrix of Fourier twiddles

```
cublas_param_mom_sum.a_stride = 0 ; // mom stays the same  
cublas_param_mom_sum.b_stride = n_spatial_sites*X[3] ;  
cublas_param_mom_sum.c_stride = X[3]*nMom ;  
cublas_param_mom_sum.batch_count = blockSizeMomProj ;
```

And we will be batching the DFT over blockSizeMomProj.

```
cublas_param_mom_sum.alpha = 1.0;  
cublas_param_mom_sum.beta = 0.0;  
cublas_param_mom_sum.data_order = QUDA_BLAS_DATAORDER_ROW;  
cublas_param_mom_sum.data_type = QUDA_BLAS_DATATYPE_Z;
```

Here we specify that the ZGEMM we want is all in row-major format

We start some copies. Here we copy all of the unbarred (inner n2) quarks to the device

```
std::vector<ColorSpinorField> quda_quark(n2) ;
for (int dil2=0; dil2<n2; dil2++) {
    quda_quark[dil2] = ColorSpinorField(cuda_quark_param) ;
    quda_quark[dil2] = quark[dil2] ;
}
cudaMemcpy(d_mom, host_mom, data_mom_bytes, cudaMemcpyHostToDevice);
```

As well as the exponentiated Fourier factors

The hot loop over n1 and n2 we have seen already in the previous lecture - the idea is to only load the barred quark inside the n1 loop. Do blockSizeMomProj traces into d_tmp (could and probably should be done in a MRHS way or streamed) and then call the batched BLAS DFT

Cleanup and copies back to the host

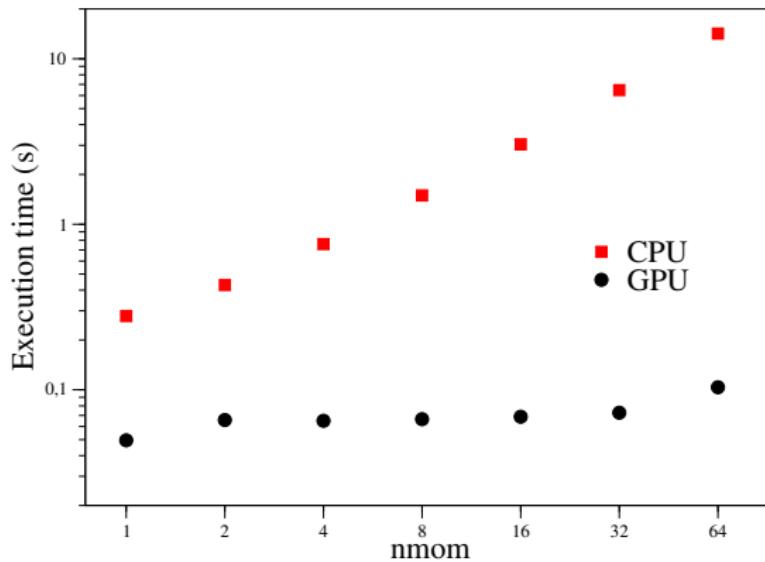
We of course must copy our object back to the host from the device

```
cudaMemcpy( return_array, d_ret, data_ret_bytes,  
            cudaMemcpyDeviceToHost );  
pool_device_free(d_ret);  
pool_device_free(d_tmp);  
pool_device_free(d_mom);
```

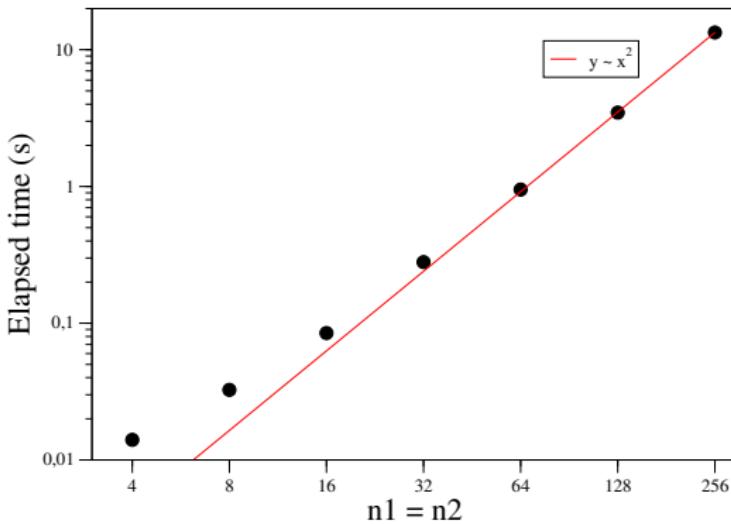
And free up the memory we have allocated

Note that we still aren't fully done yet as we have performed the DFT over the local dimensions only. It will be up to the end-user to finally put these correlators into global "t" indices. What we are returning is a small object so the burden on MPI shouldn't be bad.

Plots from Ernest



A comparison of naive single CPU vs GPU current kernels for a 16^4 local volume and $n1 = n2 = 16$ for various sizes of momenta. CPU scales linearly whereas GPU only starting to scale at large momenta.



GPU scaling with $n1 = n2$ noises, should scale like $n1^2$ and does at large $n1$ where the overhead in copies and things becomes negligible
(blockSizeMomProj=16, nMom=32)

Let's look now at the baryon kernel

```
void laphBaryonKernel( const int n1, const int n2, const int n3,
                      const double _Complex *host_coeffs1,
                      const double _Complex *host_coeffs2,
                      const double _Complex *host_coeffs3,
                      const double _Complex *host_mom,
                      const int nEv,
                      void **host_evec,
                      QudaInvertParam inv_param,
                      void *return_array,
                      const int blockSizeMomProj,
                      const int X[4] )
```

Much like the current kernel we have noise lengths $n1, n2, n3$ we also have 3 host_coeffs that are the explicit noise factors that will multiply the evecs.

```
std::vector<std::complex<double>> coeffs1(n1*nEv);
std::vector<ColorSpinorField> quda_q1(n1);
for(int i=0; i<n1; i++) {
    quda_q1[i] = ColorSpinorField(cuda_q1_param) ;
    for( int j = 0 ; j < nEv ; j++ ) {
        coeffs1[j*n1+i] = (std::complex<double>)host_coeffs1[j+i*nEv]
    }
}
```

And likewise for coeffs2 and coeffs3, this is just a cast into some format we will need in a bit. Note we are creating $n1$ evecs on the device for coeffs1 and so on for 2 and 3.

```
{  
    std::vector<ColorSpinorField> quda_evec(1) ;  
    quda_evec[0] = ColorSpinorField(cuda_evec_param);  
    for (int i=0; i<nEv; i++) {  
        quda_evec[0] = evec[i] ;  
        for( int dil1 = 0 ; dil1 < n1 ; dil1++ ) {  
            std::vector<std::complex<double>> cf(1,coeffs1[dil1+n1*i]) ;  
            blas::caxpy( cf, quda_evec, quda_q1[dil1] ) ;  
        }  
        for( int dil2 = 0 ; dil2 < n2 ; dil2++ ) {  
            std::vector<std::complex<double>> cf(1,coeffs2[dil2+n2*i]) ;  
            blas::caxpy( cf, quda_evec, quda_q2[dil2] ) ;  
        }  
        for( int dil3 = 0 ; dil3 < n3 ; dil3++ ) {  
            std::vector<std::complex<double>> cf(1,coeffs3[dil3+n3*i]) ;  
            blas::caxpy( cf, quda_evec, quda_q3[dil3] ) ;  
        }  
    }  
}
```

We create a scope where one single ColorSpinorField (device evec) will live

For every host evec we do a copy onto the device, and form the noises `quda_q*` by multiplying with the noise vector (after punning into a type that the device `blas::caxpy` wants)

We do this for every set of noises **for $n1, n2, n3 < nEv$** this is quite a big memory saving and the cost of caxpys is negligible, in principle this could be MRHS-ed if it became a problem

DFT parameters

Just like the meson setup but this time we will be ordered like $C(p, n1, n2, n3, t)$ with momenta moving slowest.

```
cublas_param_mom_sum.m = nMom;
cublas_param_mom_sum.n = X[3];
cublas_param_mom_sum.k = nSp;
cublas_param_mom_sum.lda = nSp;
cublas_param_mom_sum.ldb = nSp;
cublas_param_mom_sum.ldc = X[3]*n1*n2*n3;
cublas_param_mom_sum.a_stride = 0 ;
cublas_param_mom_sum.b_stride = nSites ;
cublas_param_mom_sum.c_stride = X[3] ;
cublas_param_mom_sum.batch_count = blockSizeMomProj;
```

Two new kernels appear!

There are two building blocks that want new QUDA kernels

`colorCrossQuda(A,B,C)`

does (over color indices)

$$C_a(x) = (A \times B)_a(x) = \sum_b \sum_c \epsilon_{abc} A_b(x) B_c(x)$$

And

`colorContractQuda(A,B,C);`

does

$$C(x) = \sum_a A_a(x) B_a(x)$$

Note that this is just like the `innerProduct` for the mesons but without a conjugate!

The hot zone

```
for( int dil1=0; dil1<n1; dil1++ ) {
    for( int dil2=0; dil2<n2; dil2++ ) {
        colorCrossQuda(quda_q1[dil1], quda_q2[dil2], quda_diq);
        for (int dil3=0; dil3<n3; dil3++) {
            colorContractQuda(quda_diq, quda_q3[dil3],
                (std::complex<double>*)d_tmp + nSites*nInBlock);
            nInBlock++;
            if (nInBlock == blockSizeMomProj ) {
                blas_lapack::native::stridedBatchGEMM(d_mom, d_tmp,
                    (std::complex<double>*)d_ret + X[3]*blockStart,
                    cublas_param_mom_sum, QUDA_CUDA_FIELD_LOCATION);
                blockStart += nInBlock ; nInBlock = 0;
            }}}}}
```

Creates a "diquark" from the color cross product of q1 and q2. Color traces with q3 and DFTs the result.

And finally the last copies and frees

```
qudaMemcpy( return_array, d_ret, data_ret_bytes,  
            qudaMemcpyDeviceToHost );  
pool_device_free(d_tmp);  
pool_device_free(d_mom);  
pool_device_free(d_ret);
```

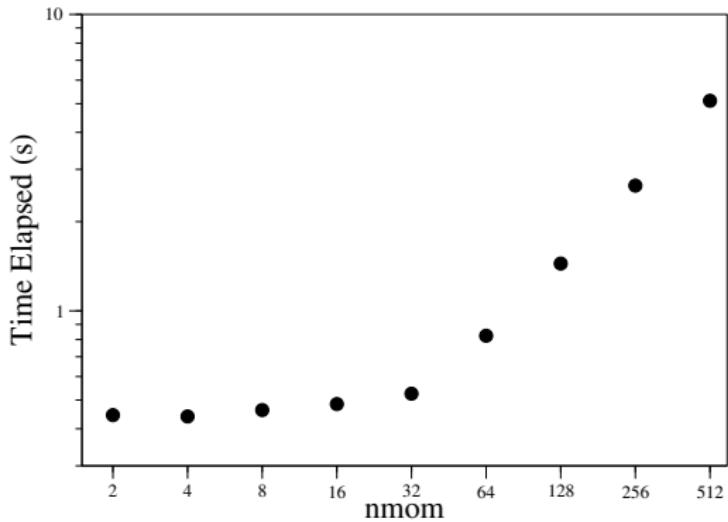
Of course the laphBaryonTripletA code works almost entirely the same way.

Homework

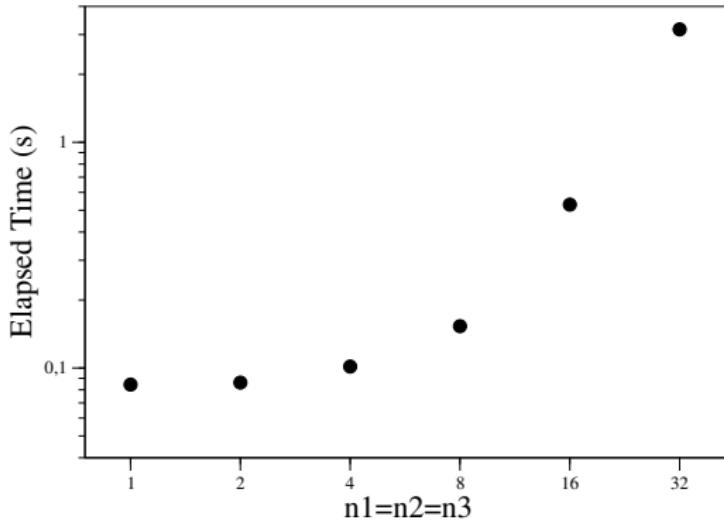
Can you implement the colorContractQuda and colorCrossQuda kernels?

Check against:

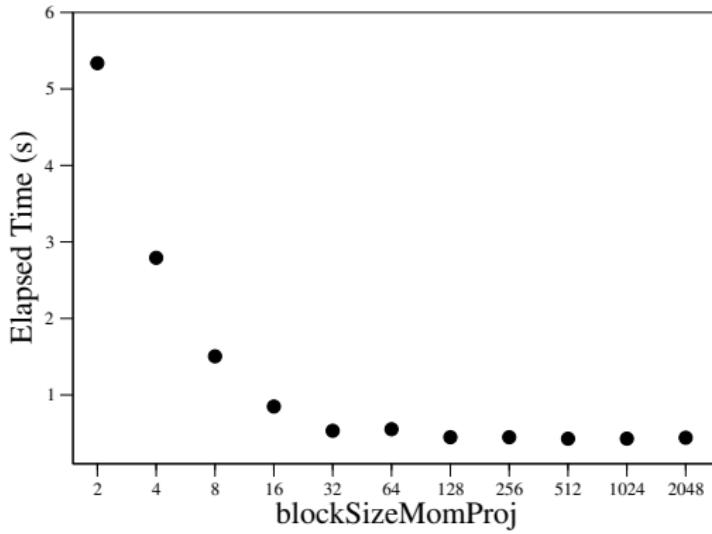
https://github.com/RJHudspith/quda/blob/develop/lib/color_contract.cu
https://github.com/RJHudspith/quda/blob/develop/include/color_spinor.h
renwick.james.hudspith@gmail.com



$\text{NeV} = 64, n1 = n2 = n3 = 16$ $\text{blockSizeMomProj}=256, 16^4$ local volume **Only for many momenta do we start to scale - DFT is not where the time is spent until 64 momenta**



$NeV = 64$, $\text{blockSizeMomProj} = n_1 * n_2$, 16^4 local volume still not close to asymptotic n^3 scaling before I ran out of memory



$\text{NeV} = 64$, $\text{nmom} = 32$, $n1 = n2 = n3 = 16$, 16^4 local volume **A factor of $10\times$ improvement from properly tuning this parameter**

A note on memory usage for the laphBaryonKernel

d_tmp is of size (double complex)*V*blockSizeMomProj

ret is of size (double complex)*n1*n2*n3*T*nmom

q1, q2, q3 are size Nc*V*(n1/n2/n3) respectively

To summarise

Current kernel with batching is fast and dominance against the CPU can be achieved by doing a large number of momenta and a large number of noises as this is where the good scaling kicks in and the overheads are better hidden.

The baryon kernel suits this as it is doing $n1 \times n2 \times n3$ traces and DFTs. For $n* < nEv$ a significant saving in memory can be achieved by just reading evects in when needed and using gratuitous caxpy calls.